

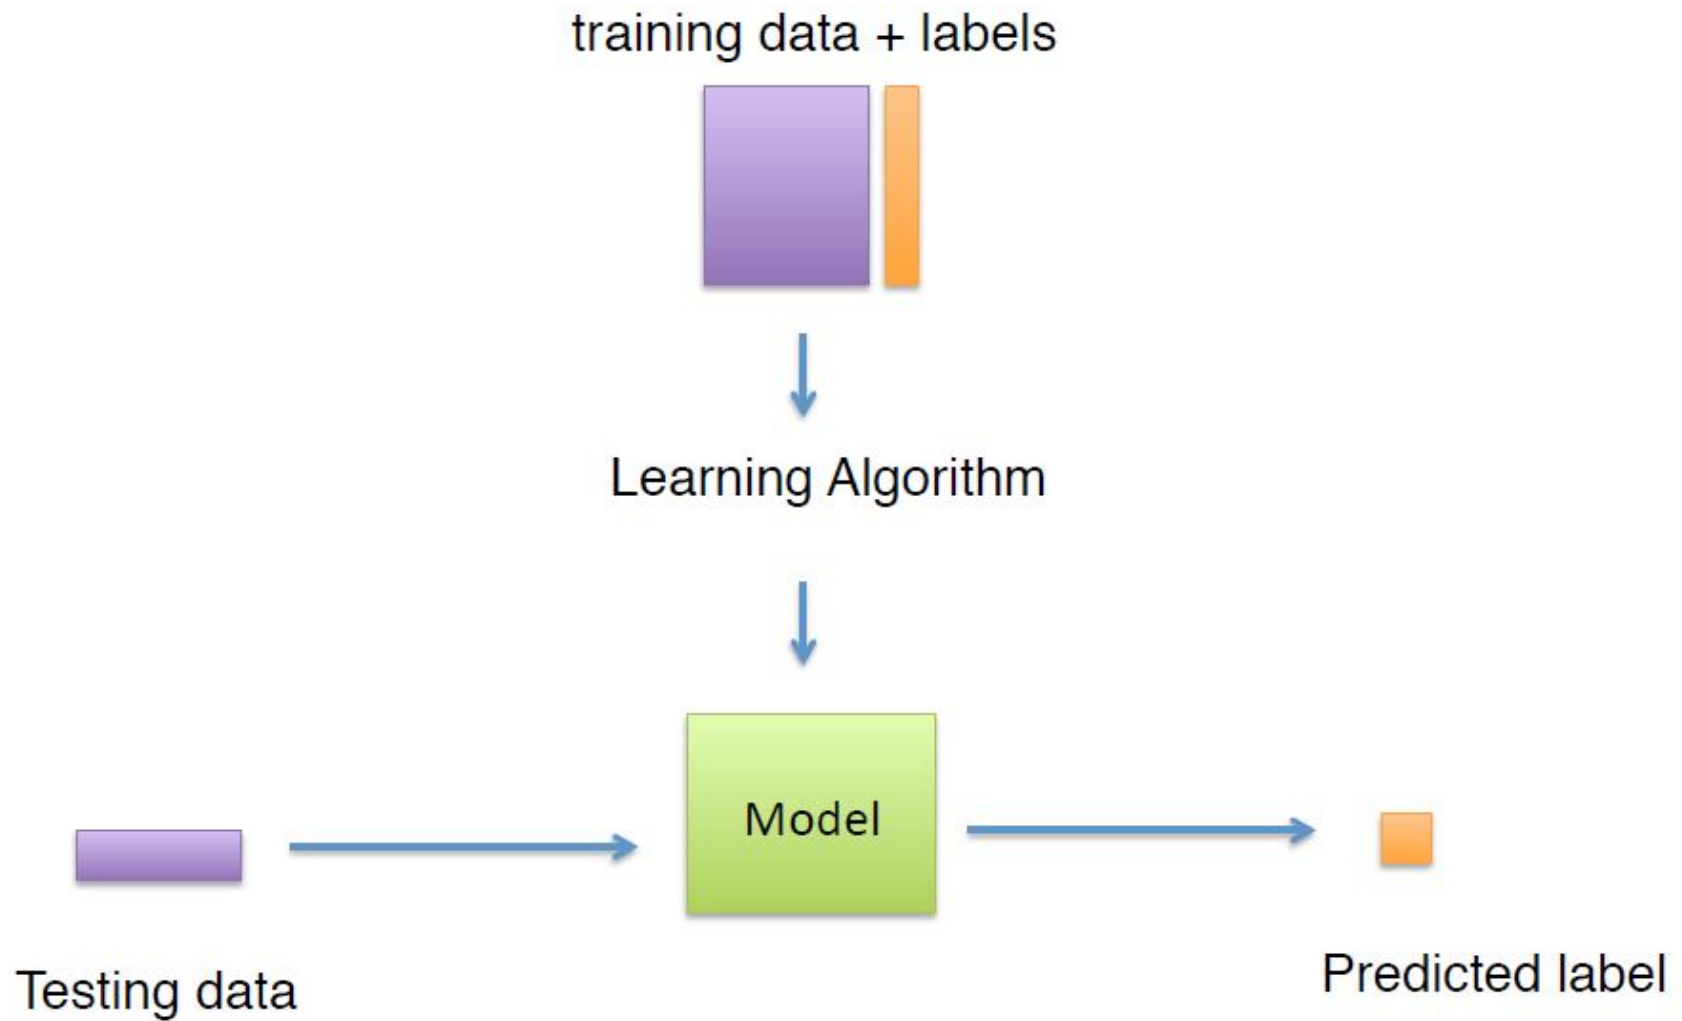
# Cost-sensitive boosting algorithms: do we really need them?

Nikos Nikolaou





The University of Manchester

# Supervised learning



# Asymmetric learning

**Cost-sensitive**  
different errors have  
have different costs




**Imbalanced classes**  
different classes appear  
with different frequency

...or both!

# Asymmetric learning



**Cost-sensitive**  
different errors have  
have different costs



**Imbalanced classes**  
different classes appear  
with different frequency

...or both!

# Boosting \ AdaBoost

- Ensemble technique: sequentially combine multiple weak learners to build a strong one

# Boosting \ AdaBoost

- Ensemble technique: sequentially combine multiple weak learners to build a strong one

## Top 10 algorithms in data mining

Xindong Wu · Vipin Kumar · J. Ross Quinlan · Joydeep Ghosh · Qiang Yang · Hiroshi Motoda · Geoffrey J. McLachlan · Angus Ng · Bing Liu · Philip S. Yu · Zhi-Hua Zhou · Michael Steinbach · David J. Hand · Dan Steinberg

Received: 9 July 2007 / Revised: 28 September 2007 / Accepted: 8 October 2007  
Published online: 4 December 2007  
© Springer-Verlag London Limited 2007

**Abstract** This paper presents the top 10 data mining algorithms identified by the IEEE International Conference on Data Mining (ICDM) in December 2006: C4.5, *k*-Means, SVM, Apriori, EM, PageRank, AdaBoost, *k*NN, Naive Bayes, and CART. These top 10 algorithms are among the most influential data mining algorithms in the research community. With each algorithm, we provide a description of the algorithm, discuss the impact of the algorithm, and review current and further research on the algorithm. These 10 algorithms cover classification,

# Boosting \ AdaBoost


- Ensemble technique: sequentially combine multiple weak learners to build a strong one

## Top 10 algorithms in data mining

Xindong Wu · Vipin Kumar · J. Ross Quinlan · Joydeep Ghosh · Qiang Yang · Hiroshi Motoda · Geoffrey J. McLachlan · Angus Ng · Bing Liu · Philip S. Yu · Zhi-Hua Zhou · Michael Steinbach · David J. Hand · Dan Steinberg

Received: 9 July 2007 / Revised: 28 September 2007 / Accepted: 8 October 2007  
Published online: 4 December 2007  
© Springer-Verlag London Limited 2007

**Abstract** This paper presents the top 10 data mining algorithms identified by the IEEE International Conference on Data Mining (ICDM) in December 2006: C4.5, *k*-Means, SVM, Apriori, EM, PageRank, AdaBoost, *k*NN, Naive Bayes, and CART. These top 10 algorithms are among the most influential data mining algorithms in the research community. With each algorithm, we provide a description of the algorithm, discuss the impact of the algorithm, and review current and further research on the algorithm. These 10 algorithms cover classification,



Face recognition in phone cameras



# Boosting \ AdaBoost

- Ensemble technique: sequentially combine multiple weak learners to build a strong one


## Top 10 algorithms in data mining

Xindong Wu · Vipin Kumar · J. Ross Quinlan · Joydeep Ghosh · Qiang Yang · Hiroshi Motoda · Geoffrey J. McLachlan · Angus Ng · Bing Liu · Philip S. Yu · Zhi-Hua Zhou · Michael Steinbach · David J. Hand · Dan Steinberg

Received: 9 July 2007 / Revised: 28 September 2007 / Accepted: 8 October 2007  
Published online: 4 December 2007  
© Springer-Verlag London Limited 2007

**Abstract** This paper presents the top 10 data mining algorithms identified by the IEEE International Conference on Data Mining (ICDM) in December 2006: C4.5, *k*-Means, SVM, Apriori, EM, PageRank, AdaBoost, *k*NN, Naive Bayes, and CART. These top 10 algorithms are among the most influential data mining algorithms in the research community. With each algorithm, we provide a description of the algorithm, discuss the impact of the algorithm, and review current and further research on the algorithm. These 10 algorithms cover classification,

kaggle™



Face recognition in phone cameras



# Boosting \ AdaBoost

- Ensemble technique: sequentially combine multiple weak learners to build a strong one

---

## An Empirical Comparison of Supervised Learning Algorithms

---

Rich Caruana  
Alexandru Niculescu-Mizil  
Department of Computer Science, Cornell University, Ithaca, NY 14853 USA

CARUANA@CS.CORNELL.EDU  
ALEXN@CS.CORNELL.EDU

With excellent performance on all eight metrics, calibrated boosted trees were the best learning algorithm overall. Random forests are close second, followed by uncalibrated bagged trees, calibrated SVMs, and uncalibrated neural nets. The models that performed


kaggle™

## Top 10 algorithms in data mining

Xindong Wu · Vipin Kumar · J. Ross Quinlan · Joydeep Ghosh · Qiang Yang · Hiroshi Motoda · Geoffrey J. McLachlan · Angus Ng · Bing Liu · Philip S. Yu · Zhi-Hua Zhou · Michael Steinbach · David J. Hand · Dan Steinberg

Received: 9 July 2007 / Revised: 28 September 2007 / Accepted: 8 October 2007  
Published online: 4 December 2007  
© Springer-Verlag London Limited 2007

**Abstract** This paper presents the top 10 data mining algorithms identified by the IEEE International Conference on Data Mining (ICDM) in December 2006: C4.5, *k*-Means, SVM, Apriori, EM, PageRank, AdaBoost, *k*NN, Naive Bayes, and CART. These top 10 algorithms are among the most influential data mining algorithms in the research community. With each algorithm, we provide a description of the algorithm, discuss the impact of the algorithm, and review current and further research on the algorithm. These 10 algorithms cover classification,



Face recognition in phone cameras

# Boosting \ AdaBoost

- Ensemble technique: sequentially combine multiple weak learners to build a strong one

---

## An Empirical Comparison of Supervised Learning Algorithms

---

Rich Caruana  
Alexandru Niculescu-Mizil  
Department of Computer Science, Cornell University, Ithaca, NY 14853 USA

CARUANA@CS.CORNELL.EDU  
ALEXN@CS.CORNELL.EDU

With excellent performance on all eight metrics, calibrated boosted trees were the best learning algorithm overall. Random forests are close second, followed by uncalibrated bagged trees, calibrated SVMs, and uncalibrated neural nets. The models that performed



2003 Gödel Prize


kaggle™

## Top 10 algorithms in data mining

Xindong Wu · Vipin Kumar · J. Ross Quinlan · Joydeep Ghosh · Qiang Yang · Hiroshi Motoda · Geoffrey J. McLachlan · Angus Ng · Bing Liu · Philip S. Yu · Zhi-Hua Zhou · Michael Steinbach · David J. Hand · Dan Steinberg


Received: 9 July 2007 / Revised: 28 September 2007 / Accepted: 8 October 2007  
Published online: 4 December 2007  
© Springer-Verlag London Limited 2007

**Abstract** This paper presents the top 10 data mining algorithms identified by the IEEE International Conference on Data Mining (ICDM) in December 2006: C4.5, *k*-Means, SVM, Apriori, EM, PageRank, AdaBoost, *k*NN, Naive Bayes, and CART. These top 10 algorithms are among the most influential data mining algorithms in the research community. With each algorithm, we provide a description of the algorithm, discuss the impact of the algorithm, and review current and further research on the algorithm. These 10 algorithms cover classification,




Face recognition in phone cameras


# AdaBoost example




# AdaBoost example




# AdaBoost example



# AdaBoost example




# AdaBoost example






# AdaBoost example



# AdaBoost example



# AdaBoost under the hood

$$\alpha_t = \frac{1}{2} \log \left( \frac{1 - \epsilon_t}{\epsilon_t} \right)$$

$$\epsilon_t = \sum_{i: h_t(\mathbf{x}_i) \neq y_i} D_i^t$$

Assign a confidence score  
to each weak learner

$$D_i^{t+1} = e^{-y_i h_t(\mathbf{x}_i) \alpha_t} D_i^t$$

Update examples' weights

$$D_i^1 = \frac{1}{N}$$

Start with a uniform weight  
distribution over the examples

$$H(\mathbf{x}') = \text{sign} \left[ \sum_{t=1}^M \alpha_t h_t(\mathbf{x}') \right]$$

Confidence weighted majority vote

# AdaBoost under the hood

$$\alpha_t = \frac{1}{2} \log \left( \frac{1 - \epsilon_t}{\epsilon_t} \right)$$
$$\epsilon_t = \sum_{i: h_t(\mathbf{x}_i) \neq y_i} D_i^t$$

Assign a confidence score  
to each weak learner

Can it handle cost-sensitive problems?

$$D_i^{t+1} = e^{-y_i h_t(\mathbf{x}_i) \alpha_t} D_i^t$$

Update examples' weights

$$D_i^1 = \frac{1}{N}$$

Start with a uniform weight  
distribution over the examples

$$H(\mathbf{x}') = \text{sign} \left[ \sum_{t=1}^M \alpha_t h_t(\mathbf{x}') \right]$$

Confidence weighted majority vote

# Asymmetric boosting variants

$$\alpha_t = \frac{1}{2} \log \left( \frac{1 - \epsilon_t}{\epsilon_t} \right)$$

$$\epsilon_t = \sum_{i: h_t(\mathbf{x}_i) \neq y_i} D_i^t$$

Assign a confidence score  
to each weak learner

$$D_i^{t+1} = e^{-y_i h_t(\mathbf{x}_i) \alpha_t} D_i^t$$

Update examples' weights

$$D_i^1 = \frac{1}{N}$$

Start with a uniform weight  
distribution over the examples

$$H(\mathbf{x}') = \text{sign} \left[ \sum_{t=1}^M \alpha_t h_t(\mathbf{x}') \right]$$

Confidence weighted majority vote

# Asymmetric boosting variants

$$\alpha_t = \frac{1}{2} \log \left( \frac{1 - \epsilon_t}{\epsilon_t} \right)$$

$$\epsilon_t = \sum_{i: h_t(\mathbf{x}_i) \neq y_i} D_i^t$$

Assign a confidence score  
to each weak learner

$$D_i^{t+1} = e^{-y_i h_t(\mathbf{x}_i) \alpha_t} D_i^t$$

Update examples' weights

(Ting & Zheng, 1998)

$$D_i^1 = \frac{1}{N}$$

Start with a uniform weight  
distribution over the examples

$$H(\mathbf{x}') = \text{sign} \left[ \sum_{t=1}^M \alpha_t h_t(\mathbf{x}') \right]$$

Confidence weighted majority vote

# Asymmetric boosting variants

(Fan et al., 1999)

$$\alpha_t = \frac{1}{2} \log \left( \frac{1 - \epsilon_t}{\epsilon_t} \right)$$
$$\epsilon_t = \sum_{i: h_t(\mathbf{x}_i) \neq y_i} D_i^t$$

Assign a confidence score  
to each weak learner

(Ting & Zheng, 1998)

$$D_i^{t+1} = e^{-y_i h_t(\mathbf{x}_i) \alpha_t} D_i^t$$

Update examples' weights

$$D_i^1 = \frac{1}{N}$$

Start with a uniform weight  
distribution over the examples

$$H(\mathbf{x}') = \text{sign} \left[ \sum_{t=1}^M \alpha_t h_t(\mathbf{x}') \right]$$

Confidence weighted majority vote



# Asymmetric boosting variants

$$\alpha_t = \frac{1}{2} \log \left( \frac{1 - \epsilon_t}{\epsilon_t} \right)$$

(Fan et al., 1999)

(Cohen & Singer, 1999)

$$\epsilon_t = \sum_{i: h_t(\mathbf{x}_i) \neq y_i} D_i^t$$

Assign a confidence score  
to each weak learner

$$D_i^{t+1} = e^{-y_i h_t(\mathbf{x}_i) \alpha_t} D_i^t$$

(Ting & Zheng, 1998)

Update examples' weights

$$D_i^1 = \frac{1}{N}$$

Start with a uniform weight  
distribution over the examples

$$H(\mathbf{x}') = \text{sign} \left[ \sum_{t=1}^M \alpha_t h_t(\mathbf{x}') \right]$$

Confidence weighted majority vote

# Asymmetric boosting variants

$$\alpha_t = \frac{1}{2} \log \left( \frac{1 - \epsilon_t}{\epsilon_t} \right)$$

$$\epsilon_t = \sum_{i: h_t(\mathbf{x}_i) \neq y_i} D_i^t$$

(Fan et al., 1999)

(Cohen & Singer, 1999)

(Ting, 2000)

Assign a confidence score  
to each weak learner

$$D_i^{t+1} = e^{-y_i h_t(\mathbf{x}_i) \alpha_t} D_i^t$$

(Ting & Zheng, 1998)

(Ting, 2000)

Update examples' weights

$$D_i^1 = \frac{1}{N}$$

Start with a uniform weight  
distribution over the examples

(Ting, 2000)

$$H(\mathbf{x}') = \text{sign} \left[ \sum_{t=1}^M \alpha_t h_t(\mathbf{x}') \right]$$

Confidence weighted majority vote

# Asymmetric boosting variants

$$\alpha_t = \frac{1}{2} \log \left( \frac{1 - \epsilon_t}{\epsilon_t} \right)$$

$$\epsilon_t = \sum_{i: h_t(\mathbf{x}_i) \neq y_i} D_i^t$$

(Fan et al., 1999)  
(Cohen & Singer, 1999)  
(Ting, 2000)  
(Joshi et al., 2001)

Assign a confidence score  
to each weak learner

$$D_i^{t+1} = e^{-y_i h_t(\mathbf{x}_i) \alpha_t} D_i^t$$

(Ting & Zheng, 1998)  
(Ting, 2000)

Update examples' weights

$$D_i^1 = \frac{1}{N}$$

Start with a uniform weight  
distribution over the examples

(Ting, 2000)

$$H(\mathbf{x}') = \text{sign} \left[ \sum_{t=1}^M \alpha_t h_t(\mathbf{x}') \right]$$

Confidence weighted majority vote

# Asymmetric boosting variants

$$\alpha_t = \frac{1}{2} \log \left( \frac{1 - \epsilon_t}{\epsilon_t} \right)$$

$$\epsilon_t = \sum_{i: h_t(\mathbf{x}_i) \neq y_i} D_i^t$$

(Fan et al., 1999)  
(Cohen & Singer, 1999)  
(Ting, 2000)  
(Joshi et al., 2001)

Assign a confidence score  
to each weak learner

$$D_i^{t+1} = e^{-y_i h_t(\mathbf{x}_i) \alpha_t} D_i^t$$

Update examples' weights

(Ting & Zheng, 1998)  
(Ting, 2000)  
(Viola & Jones, 2001; 2002)

$$D_i^1 = \frac{1}{N}$$

Start with a uniform weight  
distribution over the examples

(Ting, 2000)

$$H(\mathbf{x}') = \text{sign} \left[ \sum_{t=1}^M \alpha_t h_t(\mathbf{x}') \right]$$

Confidence weighted majority vote

# Asymmetric boosting variants

$$\alpha_t = \frac{1}{2} \log \left( \frac{1 - \epsilon_t}{\epsilon_t} \right)$$

$$\epsilon_t = \sum_{i: h_t(\mathbf{x}_i) \neq y_i} D_i^t$$

(Fan et al., 1999)

(Cohen & Singer, 1999)

(Ting, 2000)

(Joshi et al., 2001)

(Sun et al., 2005; 2007)

Assign a confidence score  
to each weak learner

$$D_i^{t+1} = e^{-y_i h_t(\mathbf{x}_i) \alpha_t} D_i^t$$

Update examples' weights

(Ting & Zheng, 1998)

(Ting, 2000)

(Viola & Jones, 2001; 2002)

$$D_i^1 = \frac{1}{N}$$

Start with a uniform weight  
distribution over the examples

(Ting, 2000)

$$H(\mathbf{x}') = \text{sign} \left[ \sum_{t=1}^M \alpha_t h_t(\mathbf{x}') \right]$$

Confidence weighted majority vote

# Asymmetric boosting variants

$$\alpha_t = \frac{1}{2} \log \left( \frac{1 - \epsilon_t}{\epsilon_t} \right)$$

$$\epsilon_t = \sum_{i: h_t(\mathbf{x}_i) \neq y_i} D_i^t$$

(Fan et al., 1999)

(Cohen & Singer, 1999)

(Ting, 2000)

(Joshi et al., 2001)

(Sun et al., 2005; 2007)

(Masnadi-Shirazi & Vasconcelos, 2007; 2011)

Assign a confidence score  
to each weak learner

$$D_i^{t+1} = e^{-y_i h_t(\mathbf{x}_i) \alpha_t} D_i^t$$

Update examples' weights

(Ting & Zheng, 1998)

(Ting, 2000)

(Viola & Jones, 2001; 2002)

$$D_i^1 = \frac{1}{N}$$

Start with a uniform weight  
distribution over the examples

(Ting, 2000)

$$H(\mathbf{x}') = \text{sign} \left[ \sum_{t=1}^M \alpha_t h_t(\mathbf{x}') \right]$$

Confidence weighted majority vote

# Asymmetric boosting variants

$$\alpha_t = \frac{1}{2} \log \left( \frac{1 - \epsilon_t}{\epsilon_t} \right)$$

$$\epsilon_t = \sum_{i: h_t(\mathbf{x}_i) \neq y_i} D_i^t$$

(Fan et al., 1999)

(Cohen & Singer, 1999)

(Ting, 2000)

(Joshi et al., 2001)

(Sun et al., 2005; 2007)

(Masnadi-Shirazi & Vasconcelos, 2007; 2011)

Assign a confidence score  
to each weak learner

$$D_i^{t+1} = e^{-y_i h_t(\mathbf{x}_i) \alpha_t} D_i^t$$

Update examples' weights

(Ting & Zheng, 1998)

(Ting, 2000)

(Viola & Jones, 2001; 2002)

$$D_i^1 = \frac{1}{N}$$

Start with a uniform weight  
distribution over the examples

(Landesa-Vázquez & Alba-Castro, 2013; 2015a; 2015b)

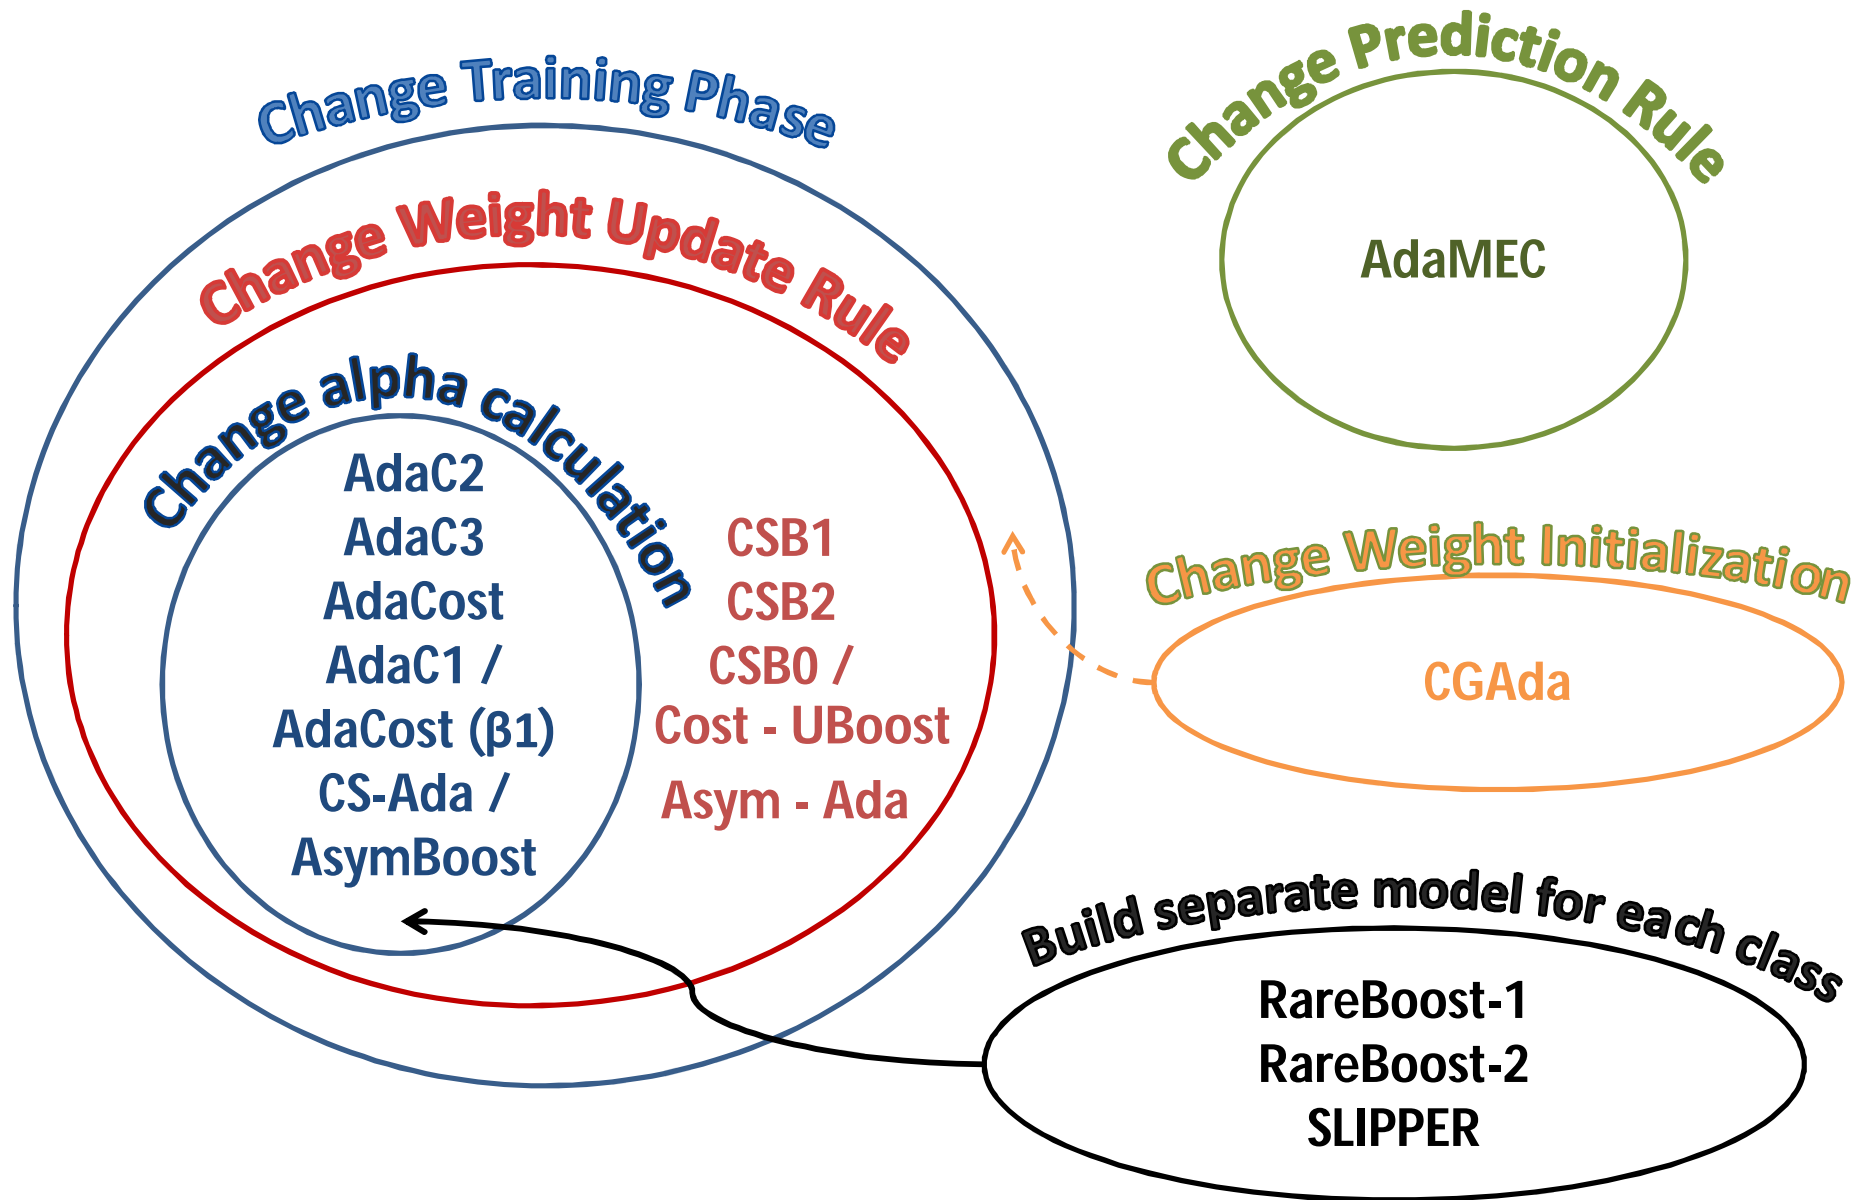
(Ting, 2000)

$$H(\mathbf{x}') = \text{sign} \left[ \sum_{t=1}^M \alpha_t h_t(\mathbf{x}') \right]$$

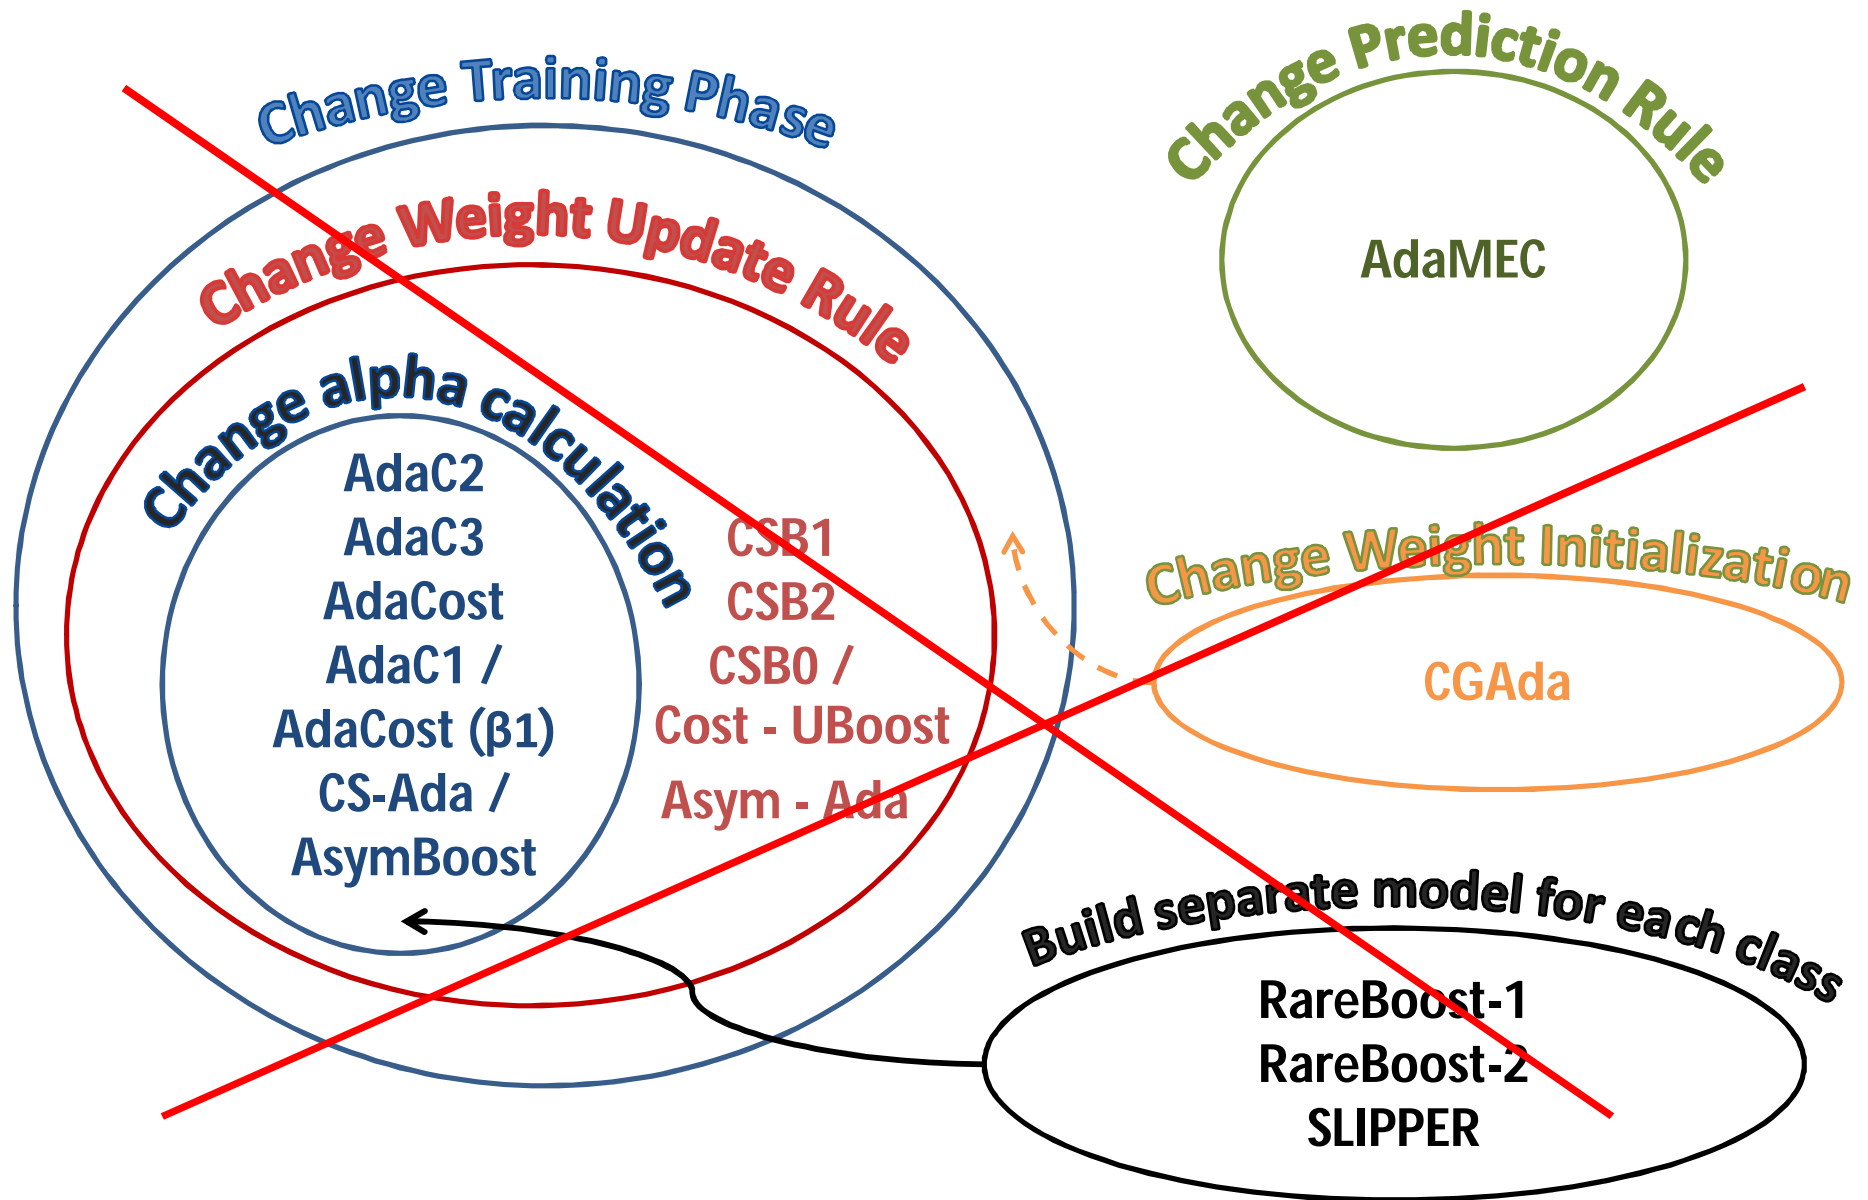
Confidence weighted majority vote



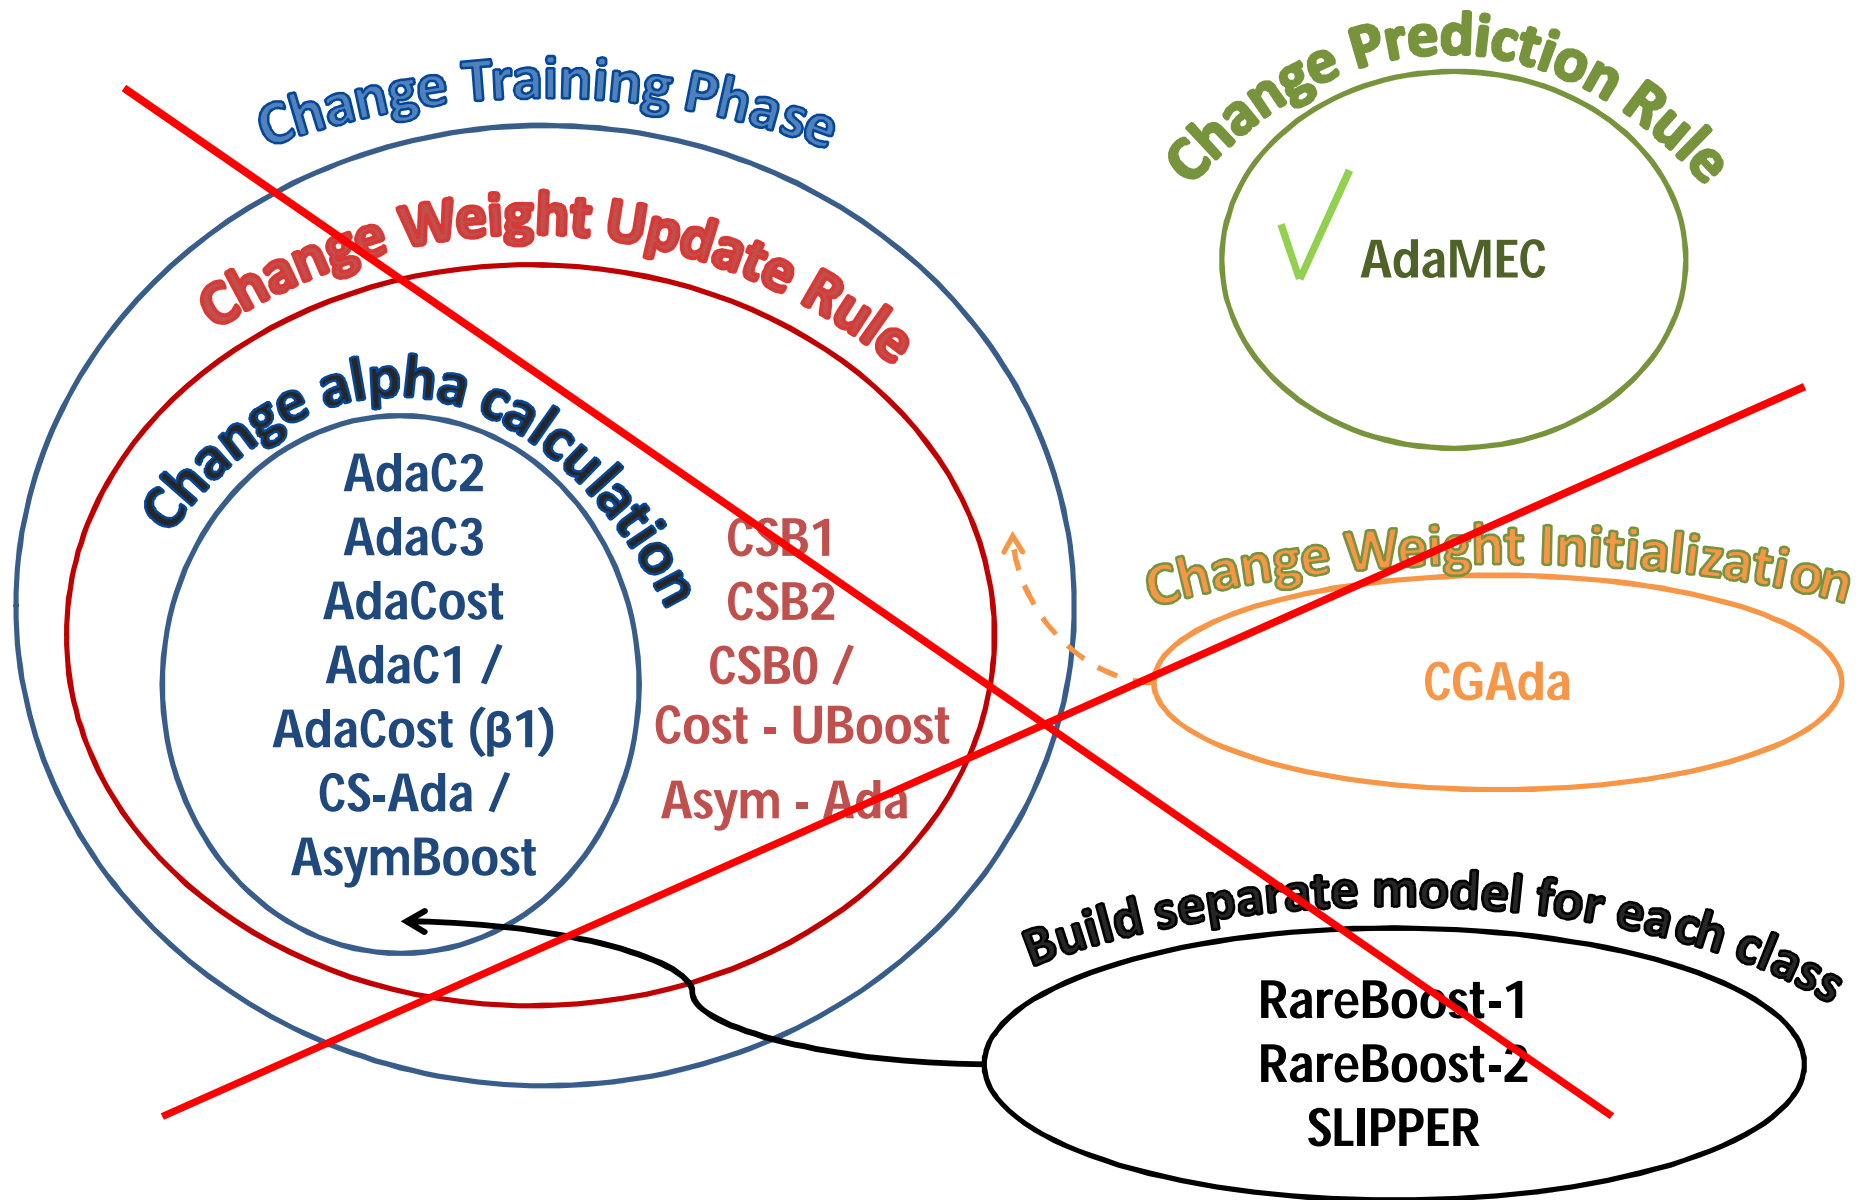
# Asymmetric boosting variants



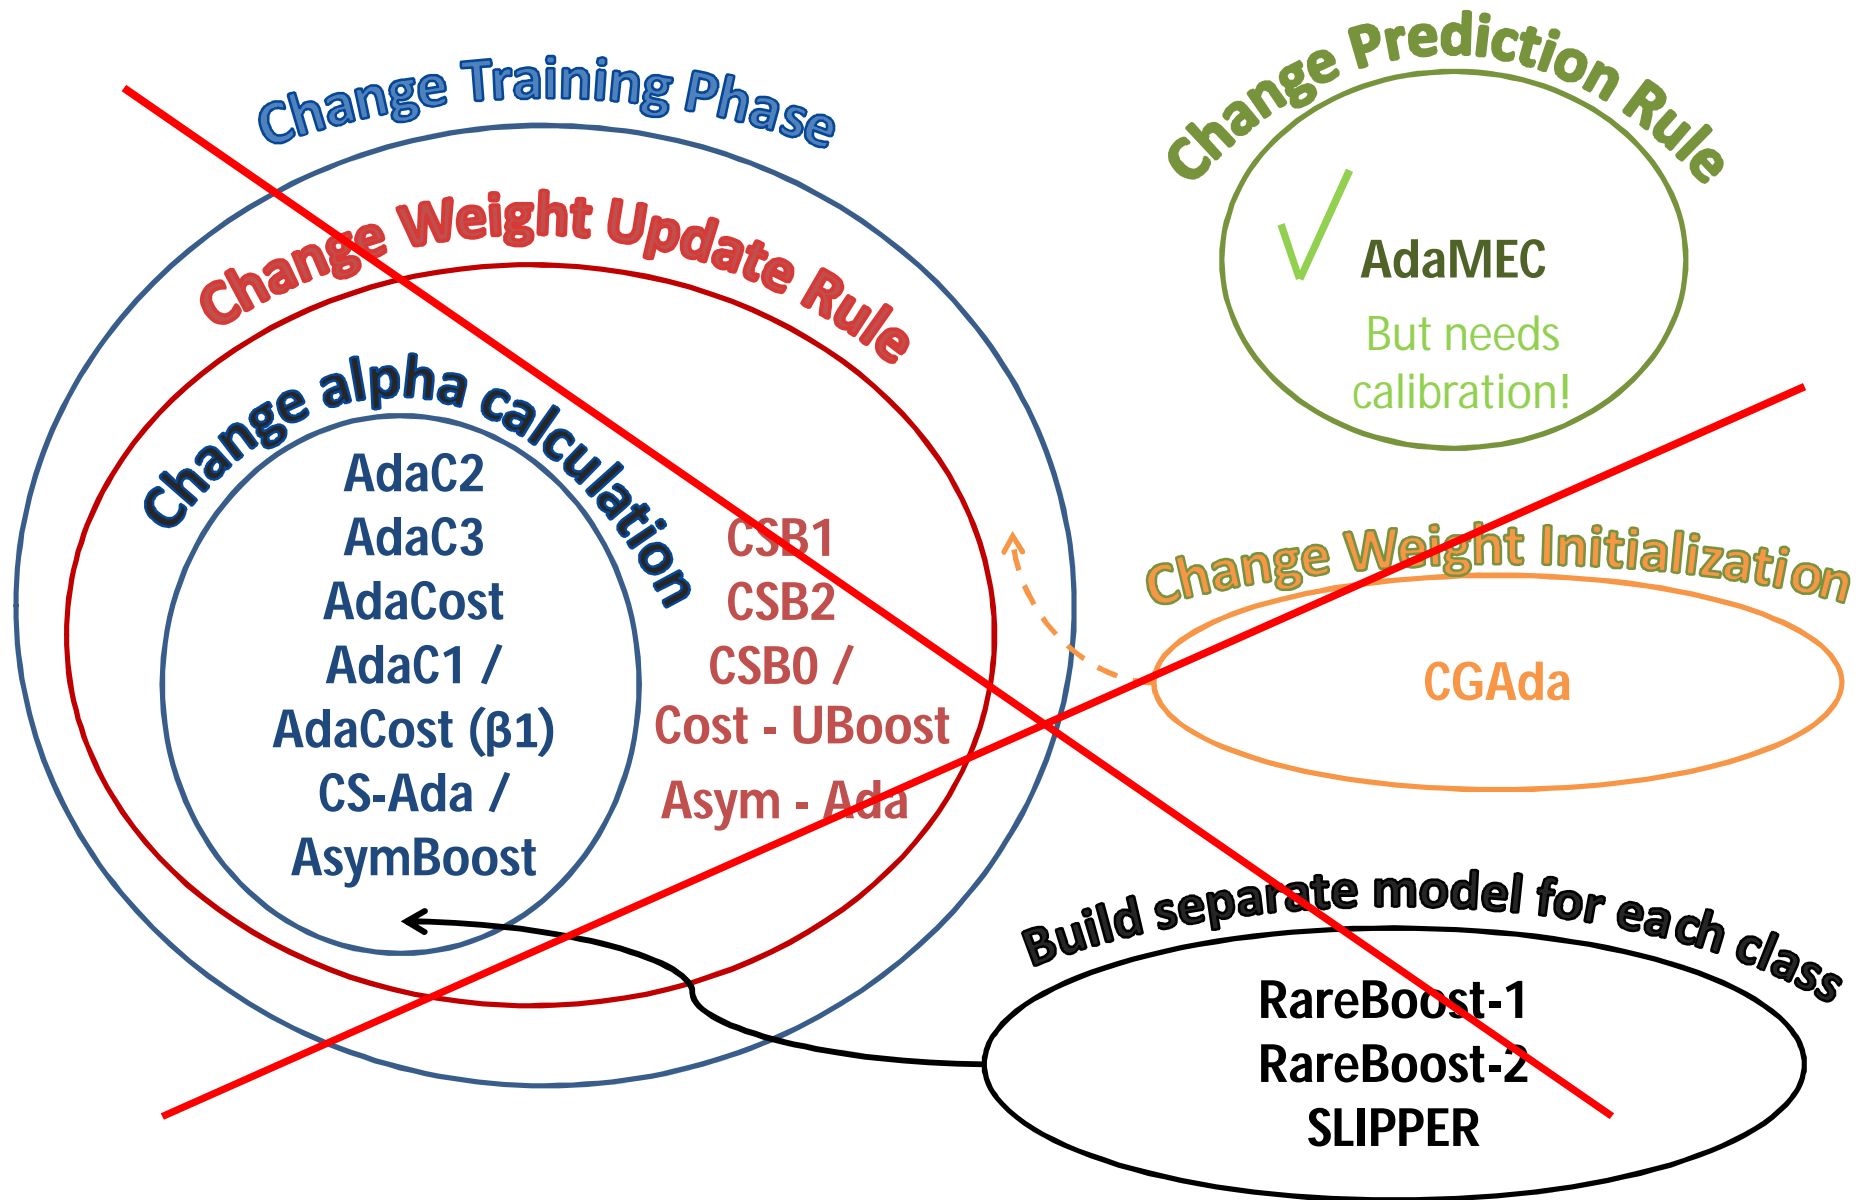
# Asymmetric boosting variants



# Asymmetric boosting variants



# Asymmetric boosting variants




# Issues with modifying training phase

- Lack **theoretical guarantees** of original AdaBoost
- Most **heuristic**, no decision-theoretic motivation
- Need to **retrain** if skew ratio changes
- Require **extra hyperparameters** to be set via CV

# Issues with *AdaMEC*

- Changes prediction rule to **minimum expected cost** ✓
- Problem: **incorrectly** assumes **scores** are **probability estimates**...
- ...but can correct this via **calibration**


# Things classifiers do...






# Things classifiers do...

- **Classify** examples
  - Is  $x$  positive?




# Things classifiers do...

- **Classify** examples
  - Is  $x$  positive?
- **Rank** examples
  - Is  $x$  'more positive' than  $x'$ ?




# Things classifiers do...

- **Classify** examples
  - Is  $x$  positive?
- **Rank** examples
  - Is  $x$  'more positive' than  $x'$ ?
- Output a **score** for each example
  - 'How positive' is  $x$ ?



# Things classifiers do...


- **Classify** examples
  - Is  $x$  positive?
- **Rank** examples
  - Is  $x$  'more positive' than  $x'$ ?
- Output a **score** for each example
  - 'How positive' is  $x$ ?
- Output a **probability estimate** for each example
  - What is the (estimated) probability that  $x$  is positive?



# Things classifiers do...

- **Classify** examples
  - Is  $x$  positive?
- **Rank** examples
  - Is  $x$  'more positive' than  $x'$ ?
- Output a **score** for each example
  - 'How positive' is  $x$ ?
- Output a **probability estimate** for each example
  - What is the (estimated) probability that  $x$  is positive?

calibration



# Why estimate probabilities?

- Need **probabilities** when a **cost-sensitive decision** needs to be made; scores won't cut it
- Will assign  $x$  to class that minimizes **expected** cost, i.e. to **Pos only** if:


expected cost of assigning  $x$  to Pos < expected cost of assigning  $x$  to Neg

$$\Leftrightarrow \hat{p}(y = 1|x) > \frac{C_{FP}}{C_{FN} + C_{FP}}$$

Costs are part of problem definition

# Probability estimates of AdaBoost

**Score for Boosting:**  $s(\mathbf{x}) = \frac{\sum_{t: h_t(x)=1} \alpha_t}{\sum_{t=1}^M \alpha_t} \in [0, 1]$




Boosted trees / stumps:  
**sigmoid distortion**;  
scores pushed more  
towards 0 or 1 as num.  
of boosting rounds  
increases

(Niculescu-Mizil & Caruana, 2006)

# Calibrating AdaBoost: Platt scaling

- Find  $A, B$  for  $\hat{p}(y = 1 | x) = \frac{1}{1 + e^{A s(x) + B}}$ , s. t. likelihood of data is maximized
- **Separate sets** for train & calibration
- Motivation: undo sigmoid distortion observed in boosted trees
- Alternative: isotonic regression





# Calibrating AdaBoost for cost-sensitive learning

On training set:

- Train AdaBoost ensemble  $H_M$



On validation set:

- Calculate score  $s(\mathbf{x}) = \frac{\sum_{t: h_t(x)=1} \alpha_t}{\sum_{t=1}^M \alpha_t} \in [0, 1]$   
of each example  $\mathbf{x}$  under ensemble  $H_M$
- Find  $A, B$  s. t. the likelihood of the data under model  $\hat{p}(y = 1|\mathbf{x}) = \frac{1}{1+e^{As(\mathbf{x})+B}}$  is maximized




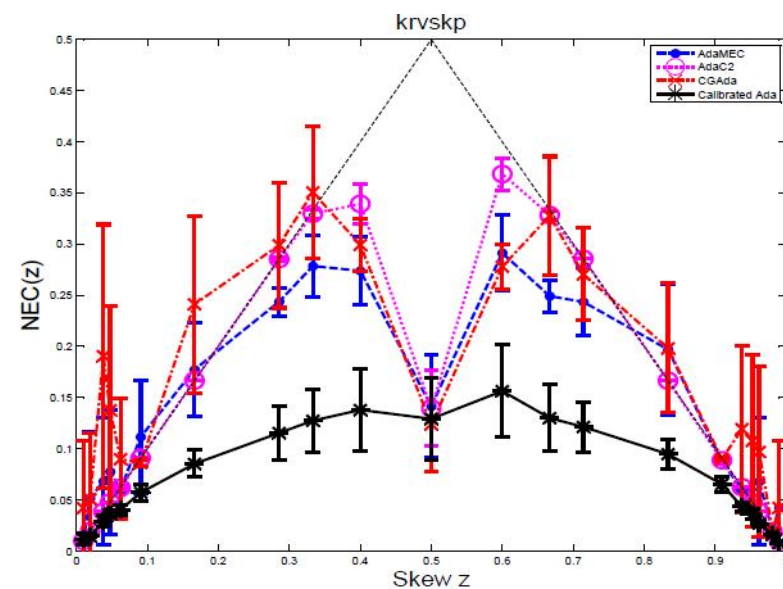


On test set:

- Calculate score  $s(\mathbf{x})$ ,  $\forall$  example  $\mathbf{x}$  under  $H_M$
- Apply transformation  $\hat{p}(y = 1|\mathbf{x}) = \frac{1}{1+e^{As(\mathbf{x})+B}}$   
to the scores  $s(\mathbf{x})$  to get probability estimates
- Predict class  $H_M(\mathbf{x}) = \text{sign} \left[ \hat{p}(y = 1|x) - \frac{C_{FP}}{C_{FP}+C_{FN}} \right]$

# Experimental design

- AdaC2 vs. CGAda vs. AdaMEC      vs.      Calibrated AdaBoost  
75% Tr / 25% Te      50% Tr / 25% Cal / 25% Te
- Weak learner: univariate logistic regression
- 18 datasets
- Evaluation: normalized expected cost  $\in [0, 1]$
- Various skew ratios:  $Z = \frac{C_{FP}}{C_{FN} + C_{FP}}$

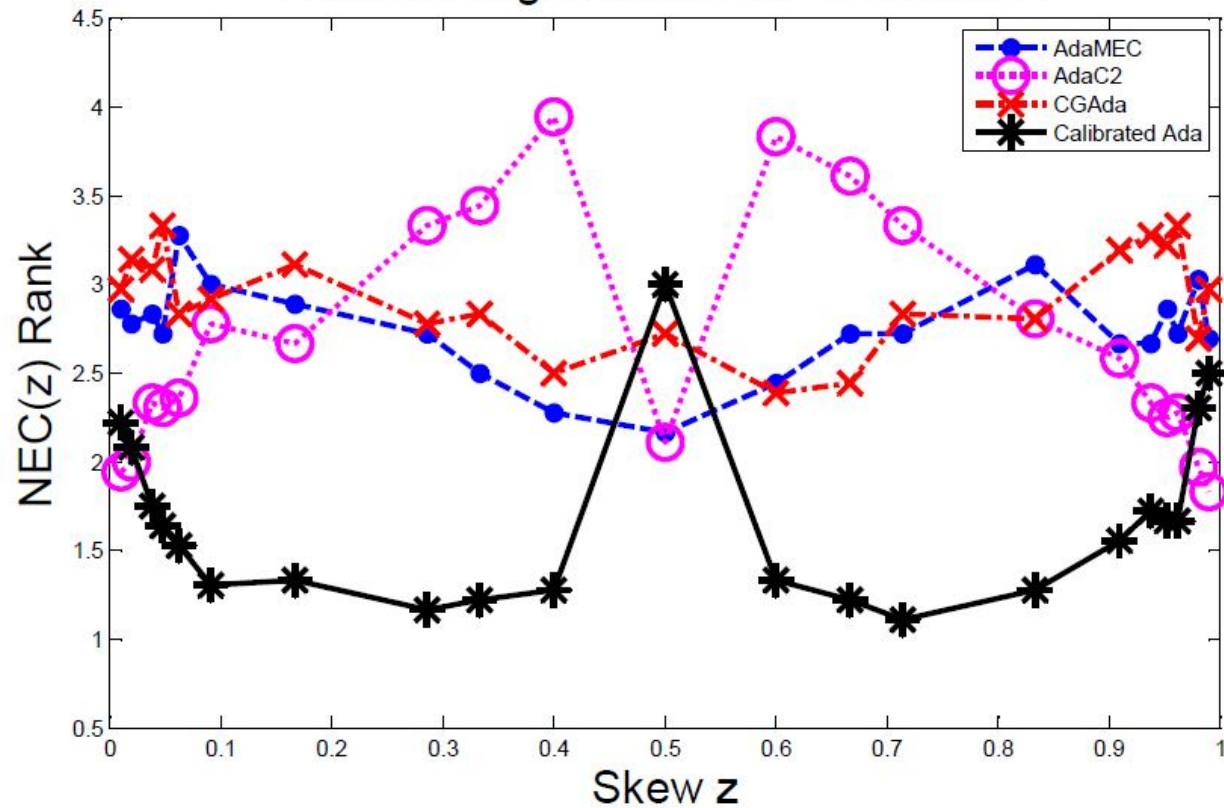
# Empirical results (1)




**Ada-Calibrated** at least as good as best, especially good on larger datasets

# Empirical results (2)

Rank averaged across all 18 datasets

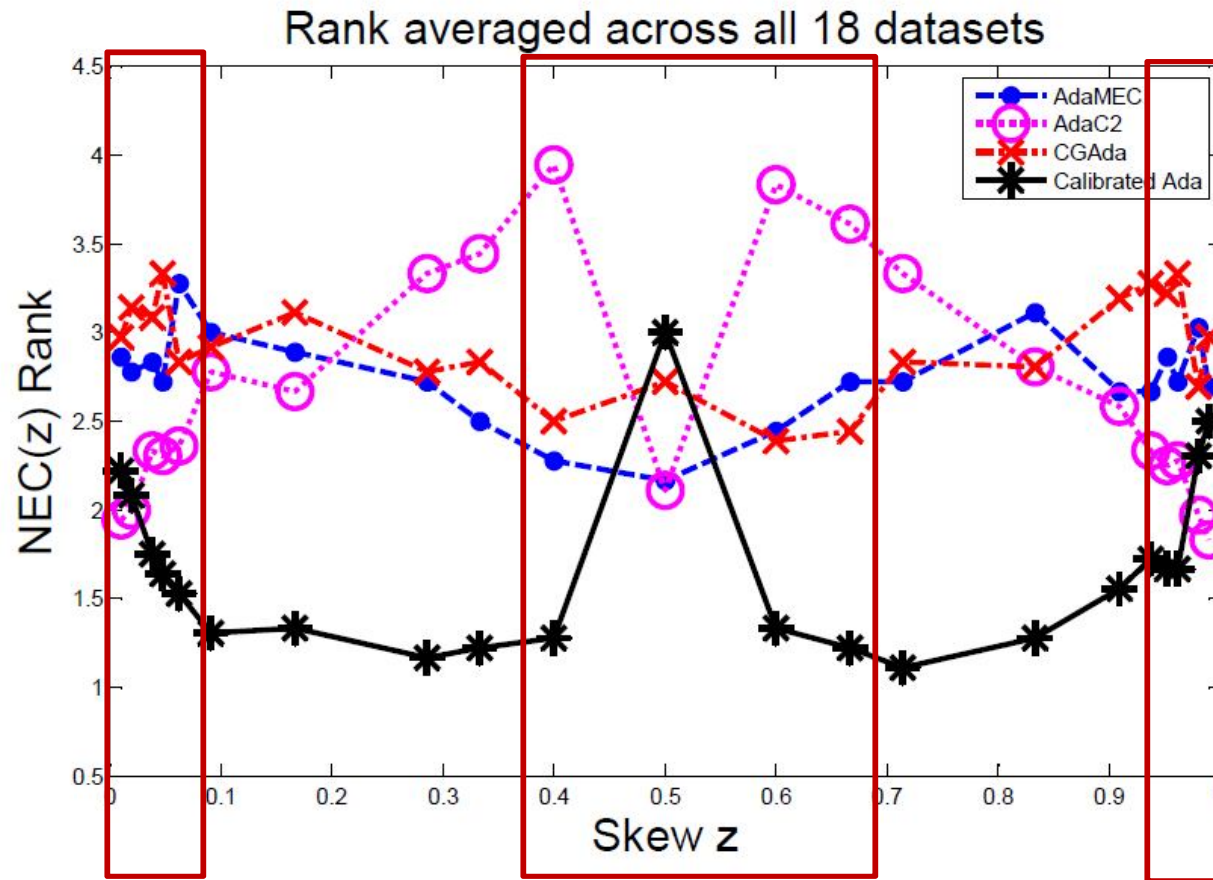


# Empirical results (2)



Nemenyi test at the 0.05 level on the differences

# Empirical results (2)




Nemenyi test at the 0.05 level on the differences

**Ada-Calibrated** at least as good as best (no sig. diff.) for very low /high skew




# Empirical results (2)



Nemenyi test at the 0.05 level on the differences

**Ada-Calibrated** at least as good as best (no sig. diff.) for very low \high skew

# Empirical results (2)



Nemenyi test at the 0.05 level on the differences

**Ada-Calibrated** at least as good as best (no sig. diff.) for very low \high skew

**Ada-Calibrated** superior to rest (sig. diff.) for medium skew



# Conclusion

- Calibrating AdaBoost empirically **at least as good as best** among alternatives published 1998 - 2015
- Conceptual **simplicity**; no need for new algorithms, or hyperparameter setting
- **No need to retrain** if skew ratio changes in deployment
- Retains **theoretical guarantees** of AdaBoosty
- Sound **probabilistic / decision-theoretic motivation**

Thank you!

# Additional Material

# Boosting as a Product of Experts

AdaBoost:  $\hat{p}(y = 1|\mathbf{x}; F_M) = \frac{\prod_{t=1}^M \hat{p}(y = 1|\mathbf{x}; f_t)}{\prod_{t=1}^M \hat{p}(y = 1|\mathbf{x}; f_t) + \prod_{t=1}^M \hat{p}(y = -1|\mathbf{x}; f_t)}$

(Edakunni et al., 2011)

# Boosting as a Product of Experts

AdaBoost:  $\hat{p}(y = 1|\mathbf{x}; F_M) = \frac{\prod_{t=1}^M \hat{p}(y = 1|\mathbf{x}; f_t)}{\prod_{t=1}^M \hat{p}(y = 1|\mathbf{x}; f_t) + \prod_{t=1}^M \hat{p}(y = -1|\mathbf{x}; f_t)}$   
(Edakunni et al., 2011)

AdaMEC:  $\hat{p}(y = 1|\mathbf{x}; F_M) = \frac{c_{FN} \prod_{t=1}^M \hat{p}(y = 1|\mathbf{x}; f_t)}{c_{FN} \prod_{t=1}^M \hat{p}(y = 1|\mathbf{x}; f_t) + c_{FP} \prod_{t=1}^M \hat{p}(y = -1|\mathbf{x}; f_t)}$

# Boosting as a Product of Experts

AdaBoost:  $\hat{p}(y = 1|\mathbf{x}; F_M) = \frac{\prod_{t=1}^M \hat{p}(y = 1|\mathbf{x}; f_t)}{\prod_{t=1}^M \hat{p}(y = 1|\mathbf{x}; f_t) + \prod_{t=1}^M \hat{p}(y = -1|\mathbf{x}; f_t)}$   
(Edakunni et al., 2011)

AdaMEC:  $\hat{p}(y = 1|\mathbf{x}; F_M) = \frac{c_{FN} \prod_{t=1}^M \hat{p}(y = 1|\mathbf{x}; f_t)}{c_{FN} \prod_{t=1}^M \hat{p}(y = 1|\mathbf{x}; f_t) + c_{FP} \prod_{t=1}^M \hat{p}(y = -1|\mathbf{x}; f_t)}$

AdaC2:  $\hat{p}(y = 1|\mathbf{x}; F_M) = \frac{c_{FN}^M \prod_{t=1}^M \hat{p}(y = 1|\mathbf{x}; f_t)}{c_{FN}^M \prod_{t=1}^M \hat{p}(y|\mathbf{x}; f_t) + c_{FP}^M \prod_{t=1}^M \hat{p}(y = -1|\mathbf{x}; f_t)}$

⋮

# Boosting as a Product of Experts

AdaBoost:  $\hat{p}(y = 1|\mathbf{x}; F_M) = \frac{\prod_{t=1}^M \hat{p}(y = 1|\mathbf{x}; f_t)}{\prod_{t=1}^M \hat{p}(y = 1|\mathbf{x}; f_t) + \prod_{t=1}^M \hat{p}(y = -1|\mathbf{x}; f_t)}$   
(Edakunni et al., 2011)

AdaMEC:  $\hat{p}(y = 1|\mathbf{x}; F_M) = \frac{c_{FN} \prod_{t=1}^M \hat{p}(y = 1|\mathbf{x}; f_t)}{c_{FN} \prod_{t=1}^M \hat{p}(y = 1|\mathbf{x}; f_t) + c_{FP} \prod_{t=1}^M \hat{p}(y = -1|\mathbf{x}; f_t)}$

AdaC2:  $\hat{p}(y = 1|\mathbf{x}; F_M) = \frac{c_{FN}^M \prod_{t=1}^M \hat{p}(y = 1|\mathbf{x}; f_t)}{c_{FN}^M \prod_{t=1}^M \hat{p}(y|\mathbf{x}; f_t) + c_{FP}^M \prod_{t=1}^M \hat{p}(y = -1|\mathbf{x}; f_t)}$

⋮

# Calibration

- $s(x) \in [0, 1]$  : score assigned by classifier to example  $x$

- A classifier is **calibrated** if

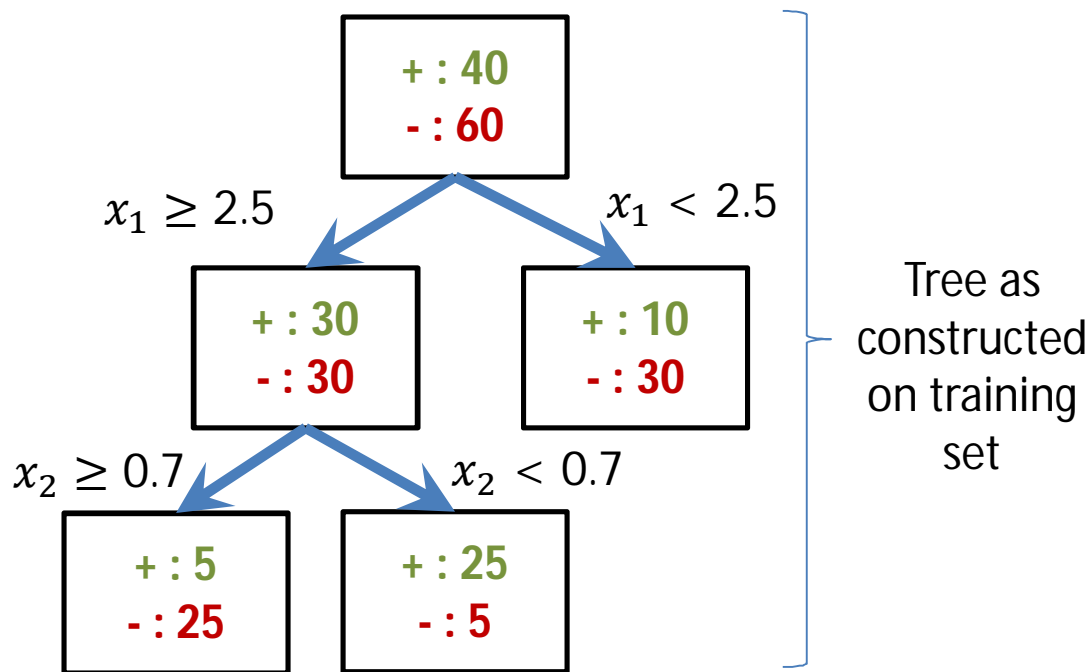
$$\hat{p}(y = 1|x) \triangleq \frac{N_{y=1,x}}{N_x} \rightarrow s(x), \text{ as } N \rightarrow \infty$$

- Intuitively: consider all examples with  $s(x) = 0.7$ ;  
70% of these examples **should** be positives
- Calibration **can only improve** classification (asymptotically)



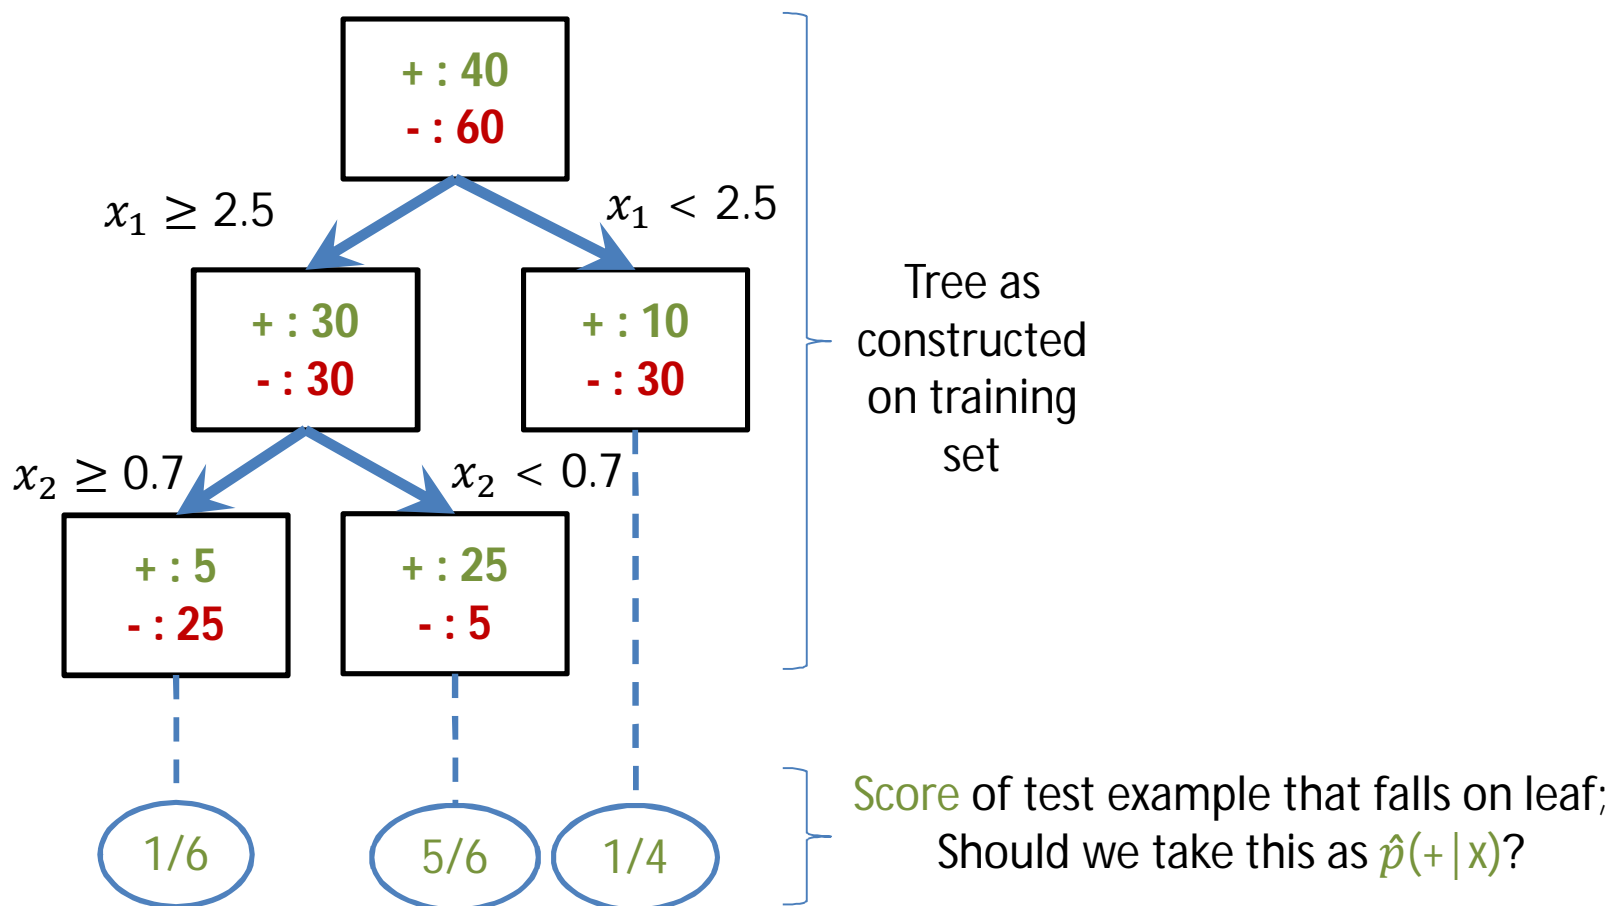
# Probability estimation is not easy

Most classifiers don't produce probability estimates **directly** but we get them via scores, e.g. decision trees:



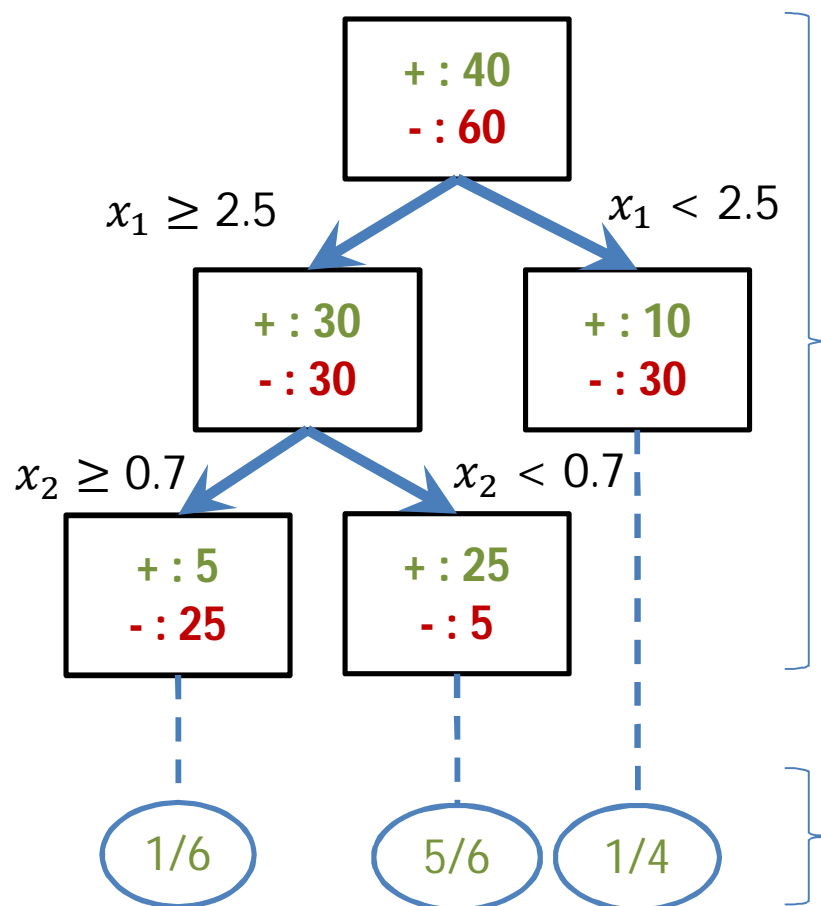
# Probability estimation is not easy

Most classifiers don't produce probability estimates **directly** but we get them via scores, e.g. decision trees:



# Probability estimation is not easy

Most classifiers don't produce probability estimates **directly** but we get them via scores, e.g. decision trees:



Tree as  
constructed  
on training  
set

Even 'probabilistic'  
classifiers can fail to  
produce **reliable**  
probability estimates  
(e.g. Naïve Bayes)

Score of test example that falls on leaf;  
Should we take this as  $\hat{p}(+|x)$ ?